

# Rappels(?) d'algorithmique

Chapter content:

- Tris
- Data structures
- Compression

# Table des matières

---

2018-2019

## Rappels(?) d'algorithmique

- Tris
- Data structures
- Compression

## Comparison sort (quizz)

Where is the snag (rather obvious)?

---

**Algorithm 1:** Sort algo in  $O(n \log n)$

---

Input:  $S = \{i_1, \dots, i_n\}$  unsorted.

Output:  $A = [o_1, \dots, o_n]$  contains  $S$ , sorted.

$A = []$

**while**  $S \neq \emptyset$

    Remove one item  $x$  from  $S$

    // no matter which:  $O(1)$

    Insert  $x$  into  $A$  at proper position

    // binary search in  $O(\log n)$

**end**

**return**  $A$

Claim: each loop in  $O(\log n)$  so  $O(n \log n)$  overall (?)

# Performance of sorts (1): comparison sorts

[https://en.wikipedia.org/wiki/Sorting\\_algorithm](https://en.wikipedia.org/wiki/Sorting_algorithm)

3 main algorithms to sort efficiently ( $n \log n$ ):

- MergeSort
- HeapSort
- QuickSort (*worst case  $O(n^2)$ , but fast on average*)

Hybrids: IntroSort (QuickSort + HeapSort), TimSort (InsertionSort + MergeSort)

Ex: gnu c++ library;

- *Introsort*: tries *quicksort*, then, if recursion exceed depth  $2 \log n$ , switches to *heapsort*.
- when input size drops below  $<16$ , insertion sort

Ex: *TimSort* (Python, Android, Java SE7, Octave);

- searches runs of consecutive items, combines small runs through insertion sort.
- applies MergeSort on sequences obtained. The merging algorithm is not trivial.

## Performance of sorts (2): sorting integers

[https://en.wikipedia.org/wiki/Sorting\\_algorithm](https://en.wikipedia.org/wiki/Sorting_algorithm)

- Input: a list of objects  $o_1, \dots, o_n$ . Each has integer weight  $1 \leq w_i \leq n$ . How would you sort?
- And if the weight is an integer between  $n$  et  $n^5$ ?

Beware:  $\Omega(n \log n)$  only holds for comparison sort.

But do not expect better in practice. Even for integers.

# Sorting Algorithms

[https://en.wikipedia.org/wiki/Sorting\\_algorithm](https://en.wikipedia.org/wiki/Sorting_algorithm)

Comparison sorts							
Name	Best	Average	Worst	Memory	Stable	Method	Other notes
Quicksort	$n \log n$ variation is $n$	$n \log n$	$n^2$	$\log n$ on average, worst case space complexity is $n$ ; Sedgwick variation is $\log n$ worst case.	Typical in-place sort is not stable; stable versions exist.	Partitioning	Quicksort is usually done in-place with $O(\log n)$ stack space. <sup>[4][5]</sup>
Merge sort	$n \log n$	$n \log n$	$n \log n$	$n$ A hybrid block merge sort is $O(1)$ mem.	Yes	Merging	Highly parallelizable (up to $O(\log n)$ using the Three Hungarians' Algorithm <sup>[6]</sup> or, more practically, Cole's parallel merge sort) for processing large amounts of data.
In-place merge sort	—	—	$n \log^2 n$ See above, for hybrid, that is $n \log n$	1	Yes	Merging	Can be implemented as a stable sort based on stable in-place merging. <sup>[7]</sup>
Heapsort	$n \log n$	$n \log n$	$n \log n$	1	No	Selection	
Insertion sort	$n$	$n^2$	$n^2$	1	Yes	Insertion	$O(n + d)$ , in the worst case over sequences that have $d$ inversions.
Introsort	$n \log n$	$n \log n$	$n \log n$	$\log n$	No	Partitioning & Selection	Used in several STL implementations.
Selection sort	$n^2$	$n^2$	$n^2$	1	No	Selection	Stable with $O(n)$ extra space, for example using lists. <sup>[8]</sup>
Timsort	$n$	$n \log n$	$n \log n$	$n$	Yes	Insertion & Merging	Makes $n$ comparisons when the data is already sorted or reverse sorted.
Cubesort	$n$	$n \log n$	$n \log n$	$n$	Yes	Insertion	Makes $n$ comparisons when the data is already sorted or reverse sorted.
Shell sort	$n \log n$	Depends on gap sequence	Depends on gap sequence; best known is $n^{4/3}$	1	No	Insertion	Small code size, no use of call stack, reasonably fast, useful where memory is at a premium such as embedded and older mainframe applications. There is a worst case $O(n (\log n)^2)$ gap sequence but it loses $O(n \log n)$ best case time.
Bubble sort	$n$	$n^2$	$n^2$	1	Yes	Exchanging	Tiny code size.
Binary tree sort	$n \log n$	$n \log n$	$n \log n$ (balanced)	$n$	Yes	Insertion	When using a self-balancing binary search tree.
Cycle sort	$n^2$	$n^2$	$n^2$	1	No	Insertion	In-place with theoretically optimal number of writes.
Library sort	$n$	$n \log n$	$n^2$	$n$	Yes	Insertion	
Patience sorting	$n$	—	$n \log n$	$n$	No	Insertion & Selection	Finds all the longest increasing subsequences in $O(n \log n)$ .
Smoothsort	$n$	$n \log n$	$n \log n$	1	No	Selection	An adaptive variant of heapsort based upon the Leonardo sequence rather than a traditional binary heap.
Strand sort	$n$	$n^2$	$n^2$	$n$	Yes	Selection	
Tournament sort	$n \log n$	$n \log n$	$n \log n$	$n^{[9]}$	No	Selection	Variation of Heap Sort.
Cocktail sort	$n$	$n^2$	$n^2$	1	Yes	Exchanging	
Comb sort	$n \log n$	$n^2$	$n^2$	1	No	Exchanging	Faster than bubble sort on average.
Gnome sort	$n$	$n^2$	$n^2$	1	Yes	Exchanging	Tiny code size.
UnShuffle Sort <sup>[10]</sup>	$n$	$k n$	$k n$	In-place for linked lists, $n \times \text{sizeof}(\text{link})$ for array.	No	Distribution and Merge	No exchanges are performed. The parameter $k$ is proportional to the entropy in the input. $k = 1$ for ordered or reverse ordered input.
Franceschini's method <sup>[11]</sup>	—	$n \log n$	$n \log n$	1	Yes	?	
Block sort	$n$	$n \log n$	$n \log n$	1	Yes	Insertion & Merging	Combine a block-based $O(n)$ in-place merge algorithm <sup>[12]</sup> with a bottom-up merge sort.
Odd-even sort	$n$	$n^2$	$n^2$	1	Yes	Exchanging	Can be run on parallel processors easily.

# Table des matières

---

2018-2019

## Rappels(?) d'algorithmique

- Tris
- Data structures
- Compression

# Data structures

[https://en.wikipedia.org/wiki/Sorting\\_algorithm](https://en.wikipedia.org/wiki/Sorting_algorithm)

Dictionary=associative array: manage key-value pairs.

(Worst-case) complexity guarantee:

hashmap

search	insert	delete
$O(1)$	$O(1)$	$O(1)$

*on average*

Maintain a set "ordered" for min/max queries?

heap (Fibonacci...)

find_min	delete_min	insert	merge
$O(1)$	$O(\log n)$	$O(1)$	$O(1)$

If we want to support predecessor/succesor/  $i^{\text{th}}$ ?

tree (AVL, RB)

search	insert	delete	pred	$i^{\text{ème}}$
$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$

or skip list, but then *on average*



# Bloom Filter

## Bloom filter:

structure de données probabiliste pour tester l'appartenance à un ensemble:

- 1 vecteur  $A$  de  $m$  bits
- $k$  fonctions de hachage  $h_1, \dots, h_k$

## Vecteur $A$ pour l'ensemble $S$ :

Au départ  $A = [0, \dots, 0]$ . Puis  
 $\forall x \in S \forall i \leq k$ , set  $A[h_i(x)]$  to 1.

## To test if $y \in S$ :

if  $(\bigwedge_i A[h_i(y)]) = 0$ :  $y \notin S$   
if  $(\bigwedge_i A[h_i(y)]) = 1$ : claim  $y \in S$   
(probably true)

Proba. of false positive:  
 $p \approx (1 - e^{-kn/m})^k$ .

## Usage:

- partition pruning (ex: joins, filters in Oracle)
- semijoins (DB2 star query optimization)

# Table des matières

---

2018-2019

## Rappels(?) d'algorithmique

- Tris
- Data structures
- Compression

## Compression: (reminders)

A dataset can be compressed efficiently if it has low entropy (= is redundant)

Huffman algorithm (or arithmetic encoding) yields optimal prefix code, but

- complex (need frequencies), often too slow
- encodes each item separately (item=symbol?)

DB applications: compress vectors. For instance, compress columns of a table (in column store). Especially useful if table is sorted on the column.

# Vector compression: widespread algorithms

- Run-length Encoding:

Ex: [a,a,a,b,b,d,e,e,e,e,e,a,a]

Compressed Vector: (a,3) (b,2) (d,1) (e,6) (a,2)

Remark: some versions record the starting offset for the run instead of length (or both).

- Delta Encoding:

Ex: [22100,22100,22101,22101,22070,22105,22105]

Compressed Vector: [22100,0,1,0,-31,35,0]

- Offset Encoding

Ex: [12,0,0,0,105,0,0,0,0,104]

Compressed Vector: (0,12) (4,105) (9,104)

- ...

Often, vector is split in blocks, which are compressed independently.

Challenge: perform operations directly on compressed data.

# Illustration: JPEG compression

[<https://en.wikipedia.org/wiki/JPEG>]

Conversion  
+ sampling

Split in blocks  
of 8x8 pixels

DCT

Quantization

RLE+Huffman

Compression

after blocks

-76	-73	-67	-62	-58	-67	-64	-55
-65	-69	-73	-38	-19	-43	-59	-56
-66	-69	-60	-15	16	-24	-62	-55
-65	-70	-57	-6	26	-22	-58	-59
-61	-67	-60	-24	-2	-40	-60	-58
-49	-63	-68	-58	-51	-60	-70	-53
-43	-57	-64	-69	-73	-67	-63	-45
-41	-49	-59	-60	-63	-52	-50	-34

after quantization

-26	3	-6	2	2	-1	0	0
0	-2	-4	1	1	0	0	0
-3	1	5	-1	-1	0	0	0
-3	1	2	-1	0	0	0	0
1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

after conversion  
3 components

$Y' C_B C_R$

after DCT

-415.38	-30.19	-61.20	27.24	56.12	-20.10	-2.39	0.46
4.47	-21.86	-60.76	10.25	13.15	-7.09	-8.54	4.88
-46.83	7.37	77.13	-24.56	-28.91	9.93	5.42	-5.65
-48.53	12.07	34.10	-14.76	-10.24	6.30	1.83	1.95
12.12	-6.55	-13.20	-3.95	-1.87	1.75	-2.79	3.14
-7.73	2.91	2.38	-5.94	-2.38	0.94	4.30	1.85
-1.03	0.18	0.42	-2.42	-0.88	-3.02	4.12	-0.66
-0.17	0.14	-1.07	-4.19	-1.17	-0.10	0.50	1.68

after RLE

(0, 2)(-3);  
(1, 2)(-3);  
(0, 1)(-2);  
(0, 2)(-6);

run length  
#bits  
value

# The Snappy library

[<https://github.com/google/snappy>]

Objective: maximize (de)compression speed, rather than compression rate.  
Devised within Google, used in many modern DBMS.

General Idea: compressed version first specifies the length of original text (as a varint), then 2 element types (same idea as for LZ77 compression)

- elements copied as is from original text
- elements recording offset and length from a repeated prefix