

# Partitioning

Chapter content:

- Partitioning

# Table of content

---

## Partitioning

- Partitioning
  - Partitioning concepts
  - Vertical partitioning
  - Horizontal partitioning: creating partitions
  - Horizontal partitioning: maintenance
  - Horizontal partitioning: benefits

# Partitioning

Principle: divide data (table, index) into pieces that can be manipulated independently.

- $\leftrightarrow$  a key feature in datawarehouses!
- several kinds of partitioning (depending on vendors too)
- choice of parameters may be crucial for performance
- transparency (hopefully): queries need not be rewritten

## Assets:

- ✓ more indexing opportunities
- ✓ easier reorganization/update (monitoring)

## Weaknesses:

- ✗ performance overhead for queries spanning multiple partitions
- ✗ more tables/complex schema

# Partitioning

Principle: divide data (table, index) into pieces that can be manipulated independently.

2 ways to partition a table:

- *vertical partitioning*: split the columns
- *horizontal partitioning*: splits the rows

We only detail partitioning in the DBMS, but also possible to partition data at the application level, or hardware.

# Vertical partitioning

## Vertical partitioning/Row splitting

splits the attribute of the table into groups. Each group is stored in a separate table.

We already saw a particular (but frequent) case of vertical partitioning:

id	price	amount	code	form	...
1	50	4	AX03F	CASH	
2	60	2	VF67D	VISA	
3	30	1	DS27W	VISA	



id	price	amount
1	50	4
2	60	2
3	30	1

id	code	form	...
1	AX03F	CASH	
2	VF67D	VISA	
3	DS27W	VISA	

# Vertical partitioning

- more tables, but fewer columns in each
- different physical devices can be used for each partition
- a view across all partitions allows to restore original rows
- used to store large (BLOBs, LONGs) or seldom-used attributes in separate table (e.g., cheaper storage media)

## Assets:

- ✓ smaller records: more can be loaded *in-memory*
- ✓ limit reorganization on frequently changing attributes
- ✓ each partition can be optimized independently
- ✓ contention (physical storage), parallelism


## Weaknesses:

- ✗ destroys semantic unit: record
- ✗ joining tables through non-materialized view degrades performance

# Horizontal partitioning: principle

## Horizontal partitioning

splits the rows of the table into groups based on attribute value. Each group is stored in a separate table.



partitions

Featured in all major DBMS: Oracle, IBM DB2, SQL Server, PostgreSQL, MySQL, MariaDB...

- allows to handle very large relations
- may simplify maintenance
- can be used to filter partitions for query optimization (partition pruning)
- allows parallelism

# Horizontal partitioning

## Horizontal partitioning

splits the rows of the table into groups based on attribute value. Each group is stored in a separate table.

id	price	amount	code	form	...
1	50	4	AX03F	CASH	
2	60	2	VF67D	VISA	
3	30	1	DS27W	VISA	
4	20	3	AP11Z	CASH	



*partitioning key*

id	price	amount	code	form	...
1	50	4	AX03F	CASH	
2	60	2	VF67D	VISA	

id	price	amount	code	form	...
3	30	1	DS27W	VISA	
4	20	3	AP11Z	CASH	



## Horizontal partitioning

splits the rows of the table into groups based on attribute value. Each group is stored in a separate table.

id	price	amount	code	form	...
1	50	4	AX03F	CASH	
2	60	2	VF67D	VISA	
3	30	1	DS27W	VISA	
4	20	3	AP11Z	CASH	



id	price	amount	code	form	...
1	50	4	AX03F	CASH	
4	20	3	AP11Z	CASH	

id	price	amount	code	form	...
3	30	1	DS27W	VISA	
2	60	2	VF67D	VISA	

## Horizontal partitioning

*... each tuple is assigned to exactly one partition depending on value of partitioning key(s).*

Oracle syntax (sketch)

```
PARTITION BY <partition_type> (attribute list) ...
```

*partitioning key*



*in multicolumn keys, column  $k + 1$  only breaks ties on columns  $(1, \dots, k)$*

# Multiple horizontal partitioning options: Oracle

## Single level partitioning

### List Partitioning

#### East Sales Region

New York  
Virginia  
Florida




#### West Sales Region

California  
Oregon  
Hawaii




#### Central Sales Region

Illinois  
Texas  
Missouri



### Range Partitioning

January and  
February



March and  
April



May and  
June



July and  
August



### Hash Partitioning




[Data Warehousing Guide, Oracle]


# Composite partitioning examples: Oracle

## Composite partitioning

**Composite Partitioning**  
Range-Hash



**Composite Partitioning**  
Range - List



...

[Data Warehousing Guide, Oracle]

# RANGE partitioning (Oracle)

PARTITION BY RANGE(column\_list) ... VALUES LESS THAN (value list)

Partitions are specified by upper bound (exclusive).

```
CREATE TABLE sales_range(sales_id int PRIMARY KEY, salesman_id NUMBER(5),
    salesman_name VARCHAR2(30), sales_amount NUMBER(10), sales_date DATE)
PARTITION BY RANGE(sales_date)
(PARTITION t21 VALUES LESS THAN(TO_DATE('02/01/2000', 'DD/MM/YYYY')),
PARTITION t22 VALUES LESS THAN(TO_DATE('03/01/2000', 'DD/MM/YYYY')),
PARTITION t23 VALUES LESS THAN(TO_DATE('04/01/2000', 'DD/MM/YYYY')),
PARTITION t24 VALUES LESS THAN(TO_DATE('05/01/2000', 'DD/MM/YYYY')));
```

MAXVALUE can be used to specify the top range, otherwise: implicit integrity constraint on the table.

[Data Warehousing Guide, Oracle]

# HASH partitioning (Oracle)

```
CREATE TABLE sales_hash  
  (salesman_id NUMBER(5), salesman_name VARCHAR2(30),  
   sales_amount NUMBER(10), week_no NUMBER(2))  
PARTITION BY HASH(salesman_id)  
PARTITIONS 4; -- number of partitions
```

For optimal distribution:

- choose adequate partitioning key  
(e.g., column(s) that are almost unique)
- take power of 2 as number of partitions (subpartitions)

## LIST partitioning (Oracle)

```
CREATE TABLE sales_list (salesman_id NUMBER(5),  
    sales_state VARCHAR2(20), sales_amount NUMBER(10))  
PARTITION BY LIST(sales_state) -- single attribute: no composite column  
    (PARTITION sales_west VALUES ('California', 'Hawaii'),  
    PARTITION sales_east VALUES ('New York', 'Virginia', 'Florida'),  
    PARTITION sales_other VALUES(DEFAULT)); -- (optional)
```



*Does not support multi-column partition keys.*

# Additional partitioning options (Oracle)

## Interval partitioning

```
CREATE TABLE sales_interv(salesman_id NUMBER(5),salesman_name VARCHAR2(30),
    sales_amount NUMBER(10),sales_date DATE)
PARTITION BY RANGE(sales_date)
INTERVAL(NUMTOYMINTERVAL(1, 'MONTH'))
(PARTITION t21 VALUES LESS THAN(TO_DATE('02/01/2000','DD/MM/YYYY')),
-- a first range partition must be declared
PARTITION t22 VALUES LESS THAN(TO_DATE('03/01/2000','DD/MM/YYYY')),
PARTITION t23 VALUES LESS THAN(TO_DATE('04/01/2000','DD/MM/YYYY')));
-- partitions after 4/01/2000 with interval width of 1 month
```

## Partition by reference

```
CREATE TABLE sales_item(sales_id int NOT NULL,
    order_id int NOT NULL, desc VARCHAR2(10),
    CONSTRAINT sales_item_fk
    FOREIGN KEY(sales_id) REFERENCES sales_range(sales_id)
)
PARTITION BY REFERENCE(sales_item_fk)
-- requires named referential constraint toward partitioned table
```

[VLDB and Partitioning Guide, Oracle]



# Additional partitioning options (Oracle)

## Partitioning on virtual column

```
CREATE TABLE sales_virt
(
  sale_id NUMBER(5)
, quantity_sold NUMBER(6) NOT NULL
, unit_price NUMBER(8,2) NOT NULL
, total_amount AS (quantity_sold*unit_price) -- cannot use PLSQL
)
PARTITION BY HASH(total_amount)
PARTITIONS 4;
```

## COMPOSITE partitioning (Oracle)

```
CREATE TABLE sales_range_hash(s_productid NUMBER, s_saledate DATE,  
    s_custid NUMBER, s_totalprice NUMBER)  
PARTITION BY RANGE (s_saledate)  
SUBPARTITION BY HASH (s_productid) SUBPARTITIONS 8  
    (PARTITION sal99q1 VALUES LESS THAN (TO_DATE('01-APR-1999', 'DD-MON-YYYY')),  
    PARTITION sal99q4 VALUES LESS THAN (TO_DATE('01-JAN-2000', 'DD-MON-YYYY')));
```

## Partitioning: which and when?

- as a general rule, consider partitioning in Oracle if:
  - table > 2Gb, more than 1M rows
  - historic data are read-only
  - multiple media storages
- RANGE partitioning: typically (not exclusively) used on temporal attributes
- HASH partitioning: to balance partition size, typical for non-temporal attributes, or to distribute data on physical storage medias
- COMPOSITE partitioning: depending on applications, if high degree of parallelism required



*unbalanced partitions tend to degrade performance*

# Horizontal partitioning: updates

ROW MOVEMENT clause

## Oracle syntax

```
CREATE TABLE ...  
  PARTITION BY ...  
  DISABLE/ENABLE ROW MOVEMENT
```

- *row movement enabled*: updating the key may cause a row to move to another partition
- *row movement disabled (default)*: update fails if it would result in row migration

# Managing partitions:maintenance

## Operations on partitions (or subpartitions)

- **ADD**: introduce a new partition (or subpartition)
- **DROP**: removes the partition and delete its content
- **TRUNCATE**: delete content of a partition
  
- **MERGE**: merges 2 partitions into a single one
- **SPLIT**: redistribute content of partition into 2 new ones
  
- **EXCHANGE**: move partition to/from table

*and also:*

MODIFY, MOVE, RENAME, COALESCE

# Add partition

ALTER TABLE <tablename> ADD PARTITION ...

```
ALTER TABLE sales_range
  ADD PARTITION jan00 VALUES LESS THAN ('01-FEB-2000')
  TABLESPACE tsx;
```

```
ALTER TABLE sales_range_hash MODIFY PARTITION sal99q1
  ADD SUBPARTITION subpart9 TABLESPACE ts99q1;
-- the rows within partition sal99q1 are re-hashed
```



*operations depend on partition type.*

# DROP partition

ALTER TABLE <tablename> DROP PARTITION ...

Fastest way to delete large volumes of data.

```
ALTER TABLE sales_range DROP PARTITION jan00;  
-- take special care for (global) indexes and integrity constraints:  
  
-- indexes: add 'UPDATE INDEXES' clause  
-- or 'ALTER INDEX sales_idx REBUILD'  
-- or first 'DELETE FROM sales_range PARTITION (jan00)' before DROP  
  
-- integrity constraints: disable them, or first DELETE
```

*Maintaining indexes during DROP may degrade performance.*

**TRUNCATE:** similar to DROP, but only rows are deleted;  
the partition (emptied) remains.



*operations depend on partition type.*

# SPLIT/MERGE partition(s)

ALTER TABLE <tablename> SPLIT PARTITION ...

```
ALTER TABLE sales_range  
  SPLIT PARTITION t21 AT (TO_DATE('01/01/2000','DD/MM/YYYY'))  
  INTO (PARTITION t21a, Partition t21b);
```

ALTER TABLE <tablename> MERGE PARTITIONS ...

```
ALTER TABLE sales_range  
  MERGE PARTITIONS t21a, t21b INTO PARTITION t21;
```



*operations depend on partition type.*



# Partition exchange

## Oracle syntax

```
ALTER TABLE <tname>  
  EXCHANGE PARTITION <pname>  
  WITH TABLE <tname2>  
  [INCLUDING/EXCLUDING INDEXES] -->  
  [WITH/WITHOUT VALIDATION] ---->
```

exchange index(es)

check rows suitable for partition

- Swaps a table and a partition.
- table and partition must have same schema, FK constraint, indexes
- extremely useful in ETL.

Typically, organize new data; indexes, constraints... in dedicated table, then integrate new data using partition exchange.

```
CREATE TABLE sales_apr_src(sales_id int PRIMARY KEY, salesman_id NUMBER(5),  
  salesman_name VARCHAR2(30), sales_amount NUMBER(10),sales_date DATE);  
  
INSERT INTO sales_apr (1,2,'T. Elliot',23,TO_DATE('04/21/2000','DD/MM/YYYY'));  
  
ALTER TABLE sales_range EXCHANGE PARTITION t24 WITH TABLE sales_apr;
```

# Partitioning benefits

(Reminder)

- allows to handle very large relations
- allows parallelism
- may simplify maintenance ✓ (ETL, archiving...)
- allows to filter partitions for query optimization (partition pruning)

## Partition pruning

Principle: *"Do not scan partitions where there can be no matching values" when the query filters values based on partitioning column.*

Optimizer selects relevant partitions, scan partition instead of whole table.

### Predicates supported:

RANGE, LIST	<i>IN, =, LIKE, range (&lt;,...</i> )
-------------	---------------------------------------

HASH	<i>IN, =</i>
------	--------------

### Oracle distinguishes

- static partition pruning: optimizer determines partition at compile time (no subquery, static predicate. . .)
- dynamic partition pruning: for statements that use subqueries. . .

## Partition-wise joins

Parallelism within a query (principle): *split the query in multiple subqueries over distinct parts of the data; process subqueries in parallel.*

Aim: speedup join processing by minimizing data exchange between parallel execution servers (reduces both communication and memory).


Partition-wise joins can be:

- *"full"* : the 2 tables are partitioned on join key
- *"partial"* : 1 of the tables must be partitioned on join key

## Full partitions-wise joins

Joined tables must be partitioned on join key, with same partition criterias (for hash: same # partitions).

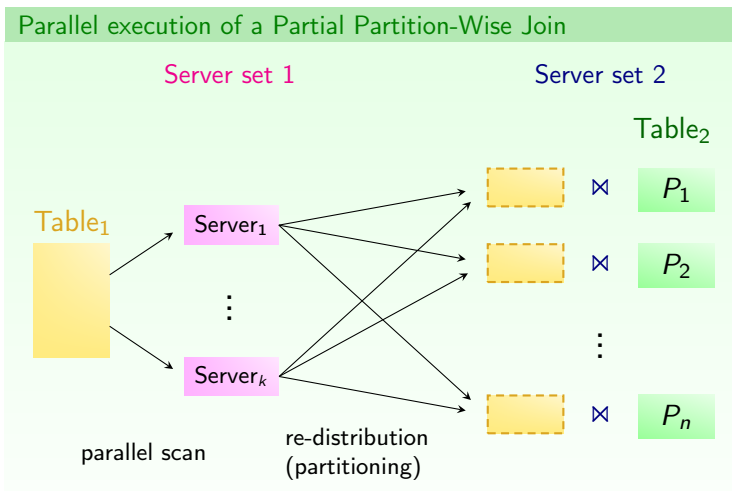
### Parallel execution of a Full Partition-Wise Join



- Can be processed serially or in parallel
- Reference partitioning is an easy way to enable full partition-wise joins.

## Partial partition-wise joins

Second table needs not be partitioned on join key, as it will be dynamically (re)partitioned based on the reference table partition.



Can be processed in parallel only.

# References

Vendors' doc:

*VLDB and Partitioning Guide*, Oracle (inspired the whole section):

<https://docs.oracle.com/database/121/VLDBG/title.htm>

<https://msdn.microsoft.com/en-us/library/ms190787.aspx>

[http://www-01.ibm.com/support/knowledgecenter/SSEPGG\\_9.7.0/com.ibm.db2.luw.admin.partition.doc/doc/c0021560.html](http://www-01.ibm.com/support/knowledgecenter/SSEPGG_9.7.0/com.ibm.db2.luw.admin.partition.doc/doc/c0021560.html)

<http://www.postgresql.org/docs/9.4/static/ddl-partitioning.html>

<https://mariadb.com/kb/en/mariadb/using-connect-partitioning-and-sharding/>