

# Physical storage in DBMS

Chapter content:

2018-2019

- Storage : DBMS adapts to hardware (Reminder)
- Reads vs Writes: storage architectures
- Techniques for handling updates
- Physical storage in RDBMS

# Table of content

---

2018-2019

## Physical storage in DBMS

- Storage : DBMS adapts to hardware (Reminder)
- Reads vs Writes: storage architectures
- Techniques for handling updates
- Physical storage in RDBMS

# New DB technologies and trends

## Data storage and processing evolutions:

- New hardware (FPGA, GPGPU, SSD, NVRAM)
- In-memory DB/ Column storage
- NoSQL: (shared nothing) massively parallel architecture (Map/Reduce, key/value)

## Trends emphasize

- supporting both *analytical* (OLAP, mining) and *prediction* queries (statistical ML)
- ACID in distributed data processing
- trust management – so far, more of an issue for transactions rather than storage/analysis (blockchain)
- integrating new (external) types of data (text, feeds, images, video...)
- building data summaries (data too large to be stored): sampling, etc.
- real-time DW

# Hardware: data storage

## The hard drive

Created in 1956 (IBM).

Multiple disks (platters) spinning rapidly together (thousands rpm/min).

Ferromagnetic. Multiple tracks per platter. 1 head per platter, on air cushion.




≈ 100€  
a few TB  
100gr.

Main storage device (secondary storage) since 1960.

The support for which DBMS were devised.

# Hardware: storage

Memory hierarchy:



On early computers, CPU frequency  $\cong$  memory bus and memory access.  
But CPU became much faster than memory.

# Table of content

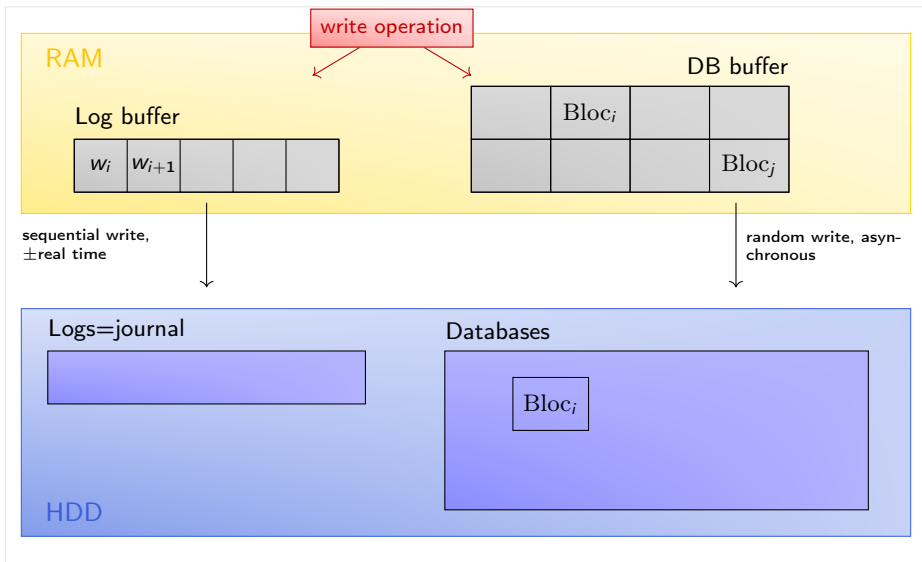
---

2018-2019


## Physical storage in DBMS

- Storage : DBMS adapts to hardware (Reminder)
- Reads vs Writes: storage architectures
- Techniques for handling updates
- Physical storage in RDBMS


# Handling writes in DBMS



# Architecture types (MVCC) in terms of tables




(a) Append-only (O2N)




(b) Append-only (N2O)



O2N : oldest-to-newest  
N2O : newest-to-oldest



(c) Time-travel Storage



(d) Delta Storage



<http://www.vldb.org/pvldb/vol10/p781-Wu.pdf> (VLDB'17)



# Table of content

---

2018-2019

## Physical storage in DBMS

- Storage : DBMS adapts to hardware (Reminder)
- Reads vs Writes: storage architectures
- **Techniques for handling updates**
- Physical storage in RDBMS

# How storage architectures deal with updates

- *update-in-place storage*: most relational DB. Limitations: versioning (SCD 2) results in fragmentation. Lock on page during update. Strategy often used with B-tree index. Provides optimal reads ( $\approx 1$  seek per read/1 per scan if unfragmented). But updating index slower.
- *log-structured storage*: all updates appended as a log (no “main”). Limitations: scans get expensive: must read all logs. But writes faster than in B-tree (and no lock). A popular version: *LSM-tree* (Log-Structured Merge tree).
- *delta with main store*: main store is read-optimized, updates recorded in write-optimized buffer (*SAP Hana/Sans Souci, and other column stores*). Buffer merged from time to time.

⇒ **Rationale for log-structured storage: instead of writing a full page on each update, defer and process them in batch. Idea of *deferred writes* common with delta approach.**

# LSM-tree

Storage optimized for write-throughput (as every log-structured storage):

- LSM-tree consists of 2 or more layers  $C_0, C_1, C_2, \dots, C_k$ .
- $C_0$  in-memory, can be tree, hash table...
- others ( $C_1, \dots$ ) are B-trees, recorded on disk/slower memory.
- updates written only in  $C_0$
- $\forall i$  when  $C_i$  is full, we merge it into  $C_{i+1}$  ( $C_i$  is emptied)
  - $\hookrightarrow \forall i$  versions in  $C_i$  always more recent than in  $C_{i+1}$
  - $\hookrightarrow$  at most  $r$  versions of record co-exist.
  - $\hookrightarrow$  a Read must search layers  $C_0, C_1, \dots$  until item found.

...but we did not detail what "full" means:

Original version:

[O'Neill et al. 1996]

- recommend using geometric progression:  $|C_{i+1}| = r \times |C_i|$ 
  - $\hookrightarrow$  total number of unique keys:  $N = O(r^k)$ .
- but claimed inefficient (does not exploit write locality)

# LSM-tree: assets and drawbacks (compared to Btree)

Buffered writes makes writes faster.

## *Pros:*

- ✓ write operations are faster (sequential)
- ✓ lower fragmentation (so range queries get more efficient)

## *Weaknessess:*

- ✗ accessing a given value may be slow



mongoDB

APACHE  
HBASE



LEVELDB

# References

## Rappels simples sur le stockage dans les SGBDs

- [\*http://sys.bdpedia.fr/\*](http://sys.bdpedia.fr/)

## MVCC storage

- [\*http://www.vldb.org/pvldb/vol10/p781-Wu.pdf\*](http://www.vldb.org/pvldb/vol10/p781-Wu.pdf)

## LSM trees

- [\*http://www.vldb.org/pvldb/vol10/p1526-bocksrocker.pdf\*](http://www.vldb.org/pvldb/vol10/p1526-bocksrocker.pdf)
- [\*http://www.eecs.harvard.edu/~margo/cs165/papers/gp-lsm.pdf\*](http://www.eecs.harvard.edu/~margo/cs165/papers/gp-lsm.pdf)
- [\*https://www.quora.com/How-does-the-Log-Structured-Merge-Tree-work\*](https://www.quora.com/How-does-the-Log-Structured-Merge-Tree-work)

# Table of content

---

2018-2019

## Physical storage in DBMS

- Storage : DBMS adapts to hardware (Reminder)
- Reads vs Writes: storage architectures
- Techniques for handling updates
- Physical storage in RDBMS

## 3 levels of design

Conceptual design      User-oriented description, independent of implementation

*E-R, UML*


Logical design      Logical description, independent of DBMS

*Relational model: table schema*

Physical design      Actual database structures

*materialized views, partitions, indexes*

# Accessing data



- : logical schema (ANSI SPARC architecture)
- : physical schema (ANSI SPARC architecture)




## Physical storage in DBMS (reminder)


- physical address of record provided by ROWID
- Records stored in pages=blocks.
- Read/Write operations require the corresponding data be brought in buffer (main memory).
- DBMS must minimize pages I/O for performance.

larger block size: more tuples edited per page I/O, but fewer pages in main-memory.

# Physical storage: Oracle (1)




**E-R diagram of  
storage structures**



**logical storage**

[Oracle Database Concepts]

## Physical storage: Oracle (2)



# References

*Oracle Database Concepts*

<https://docs.oracle.com/database/121/CNCPT/logical.htm>

[https://en.wikibooks.org/wiki/Oracle\\_and\\_DB2,\\_Comparison\\_and\\_Compatibility/Storage\\_Model/Physical\\_Storage/Summary](https://en.wikibooks.org/wiki/Oracle_and_DB2,_Comparison_and_Compatibility/Storage_Model/Physical_Storage/Summary)

ftp:  
[//ftp.software.ibm.com/software/data/db2/9/labchats/20110331-slides.pdf](ftp://ftp.software.ibm.com/software/data/db2/9/labchats/20110331-slides.pdf)  
(a bit outdated: 2011...)

<http://www.postgresql.org/docs/9.4/static/storage-page-layout.html>